

```

import pandas as pd

# Load training data for FD001
train_fd001 = pd.read_csv('train_FD001.txt', sep=' ', header=None)

# Load test data for FD001
test_fd001 = pd.read_csv('test_FD001.txt', sep=' ', header=None)

# Load true RUL values for FD001
rul_fd001 = pd.read_csv('RUL_FD001.txt', sep=' ', header=None)

train_fd001 = train_fd001.dropna(axis=1, how='all')
test_fd001 = test_fd001.dropna(axis=1, how='all')

column_names = [
    'unit_number', 'time_in_cycles', 'operational_setting_1',
    'operational_setting_2',
    'operational_setting_3'] + [f'sensor_measurement_{i}' for i in
range(1, 22)]

train_fd001.columns = column_names
test_fd001.columns = column_names

print(train_fd001.head())
print(test_fd001.head())
print(rul_fd001.head())

```

| | unit_number | time_in_cycles | operational_setting_1 | |
|-------------------------|-------------|----------------|-----------------------|---|
| operational_setting_2 \ | | | | |
| 0 | 1 | 1 | -0.0007 | - |
| 0.0004 | | | | |
| 1 | 1 | 2 | 0.0019 | - |
| 0.0003 | | | | |
| 2 | 1 | 3 | -0.0043 | |
| 0.0003 | | | | |
| 3 | 1 | 4 | 0.0007 | |
| 0.0000 | | | | |
| 4 | 1 | 5 | -0.0019 | - |
| 0.0002 | | | | |

| | operational_setting_3 | sensor_measurement_1 | |
|------------------------|-----------------------|----------------------|--------|
| sensor_measurement_2 \ | | | |
| 0 | 100.0 | 518.67 | 641.82 |
| 1 | 100.0 | 518.67 | 642.15 |
| 2 | 100.0 | 518.67 | 642.35 |
| 3 | 100.0 | 518.67 | 642.35 |
| 4 | 100.0 | 518.67 | 642.37 |

| | sensor_measurement_3 | sensor_measurement_4 |
|-----------|----------------------|----------------------|
| 0 | 1589.70 | 1400.60 |
| 14.62 ... | | |
| 1 | 1591.82 | 1403.14 |
| 14.62 ... | | |
| 2 | 1587.99 | 1404.20 |
| 14.62 ... | | |
| 3 | 1582.79 | 1401.87 |
| 14.62 ... | | |
| 4 | 1582.85 | 1406.22 |
| 14.62 ... | | |

| | sensor_measurement_12 | sensor_measurement_13 | sensor_measurement_14 |
|---|-----------------------|-----------------------|-----------------------|
| 0 | 521.66 | 2388.02 | 8138.62 |
| 1 | 522.28 | 2388.07 | 8131.49 |
| 2 | 522.42 | 2388.03 | 8133.23 |
| 3 | 522.86 | 2388.08 | 8133.83 |
| 4 | 522.19 | 2388.04 | 8133.80 |

| | sensor_measurement_15 | sensor_measurement_16 | sensor_measurement_17 |
|---|-----------------------|-----------------------|-----------------------|
| 0 | 8.4195 | 0.03 | 392 |
| 1 | 8.4318 | 0.03 | 392 |
| 2 | 8.4178 | 0.03 | 390 |
| 3 | 8.3682 | 0.03 | 392 |
| 4 | 8.4294 | 0.03 | 393 |

| | sensor_measurement_18 | sensor_measurement_19 | sensor_measurement_20 |
|---|-----------------------|-----------------------|-----------------------|
| 0 | 2388 | 100.0 | 39.06 |
| 1 | 2388 | 100.0 | 39.00 |
| 2 | 2388 | 100.0 | 38.95 |
| 3 | 2388 | 100.0 | 38.88 |

| | | | |
|---|------|-------|-------|
| 4 | 2388 | 100.0 | 38.90 |
|---|------|-------|-------|

| | |
|-----------------------|---------|
| sensor_measurement_21 | |
| 0 | 23.4190 |
| 1 | 23.4236 |
| 2 | 23.3442 |
| 3 | 23.3739 |
| 4 | 23.4044 |

[5 rows x 26 columns]

| | | | |
|-------------|----------------|-----------------------|-------------------------|
| unit_number | time_in_cycles | operational_setting_1 | operational_setting_2 \ |
| 0 | 1 | 1 | 0.0023 |
| 0.0003 | | | |
| 1 | 1 | 2 | -0.0027 |
| 0.0003 | | | |
| 2 | 1 | 3 | 0.0003 |
| 0.0001 | | | |
| 3 | 1 | 4 | 0.0042 |
| 0.0000 | | | |
| 4 | 1 | 5 | 0.0014 |
| 0.0000 | | | |

| | |
|------------------------|----------------------|
| operational_setting_3 | sensor_measurement_1 |
| sensor_measurement_2 \ | |
| 0 | 100.0 |
| | 518.67 |
| | 643.02 |
| 1 | 100.0 |
| | 518.67 |
| | 641.71 |
| 2 | 100.0 |
| | 518.67 |
| | 642.46 |
| 3 | 100.0 |
| | 518.67 |
| | 642.44 |
| 4 | 100.0 |
| | 518.67 |
| | 642.51 |

| | |
|----------------------------|----------------------|
| sensor_measurement_3 | sensor_measurement_4 |
| sensor_measurement_5 ... \ | |
| 0 | 1585.29 |
| | 1398.21 |
| 14.62 ... | |
| 1 | 1588.45 |
| | 1395.42 |
| 14.62 ... | |
| 2 | 1586.94 |
| | 1401.34 |
| 14.62 ... | |
| 3 | 1584.12 |
| | 1406.42 |
| 14.62 ... | |
| 4 | 1587.19 |
| | 1401.92 |
| 14.62 ... | |

| | sensor_measurement_12 | sensor_measurement_13 | sensor_measurement_14 |
|-----------------------|-----------------------|-----------------------|-----------------------|
| \ | | | |
| 0 | 521.72 | 2388.03 | 8125.55 |
| 1 | 522.16 | 2388.06 | 8139.62 |
| 2 | 521.97 | 2388.03 | 8130.10 |
| 3 | 521.38 | 2388.05 | 8132.90 |
| 4 | 522.15 | 2388.03 | 8129.54 |
| | | | |
| | sensor_measurement_15 | sensor_measurement_16 | sensor_measurement_17 |
| \ | | | |
| 0 | 8.4052 | 0.03 | 392 |
| 1 | 8.3803 | 0.03 | 393 |
| 2 | 8.4441 | 0.03 | 393 |
| 3 | 8.3917 | 0.03 | 391 |
| 4 | 8.4031 | 0.03 | 390 |
| | | | |
| | sensor_measurement_18 | sensor_measurement_19 | sensor_measurement_20 |
| \ | | | |
| 0 | 2388 | 100.0 | 38.86 |
| 1 | 2388 | 100.0 | 39.02 |
| 2 | 2388 | 100.0 | 39.08 |
| 3 | 2388 | 100.0 | 39.00 |
| 4 | 2388 | 100.0 | 38.99 |
| | | | |
| | sensor_measurement_21 | | |
| 0 | 23.3735 | | |
| 1 | 23.3916 | | |
| 2 | 23.4166 | | |
| 3 | 23.3737 | | |
| 4 | 23.4130 | | |
| [5 rows x 26 columns] | | | |
| | 0 1 | | |
| 0 | 112 NaN | | |
| 1 | 98 NaN | | |
| 2 | 69 NaN | | |

```
3    82 NaN
4    91 NaN
```

```
train_fd001.to_csv('train_fd001.csv', index=False)
test_fd001.to_csv('test_fd001.csv', index=False)
rul_fd001.to_csv('rul_fd001.csv', index=False)
```

```
print(train_fd001.isnull().sum())
```

```
unit_number          0
time_in_cycles       0
operational_setting_1 0
operational_setting_2 0
operational_setting_3 0
sensor_measurement_1  0
sensor_measurement_2  0
sensor_measurement_3  0
sensor_measurement_4  0
sensor_measurement_5  0
sensor_measurement_6  0
sensor_measurement_7  0
sensor_measurement_8  0
sensor_measurement_9  0
sensor_measurement_10 0
sensor_measurement_11 0
sensor_measurement_12 0
sensor_measurement_13 0
sensor_measurement_14 0
sensor_measurement_15 0
sensor_measurement_16 0
sensor_measurement_17 0
sensor_measurement_18 0
sensor_measurement_19 0
sensor_measurement_20 0
sensor_measurement_21 0
dtype: int64
```

```
import matplotlib.pyplot as plt
```

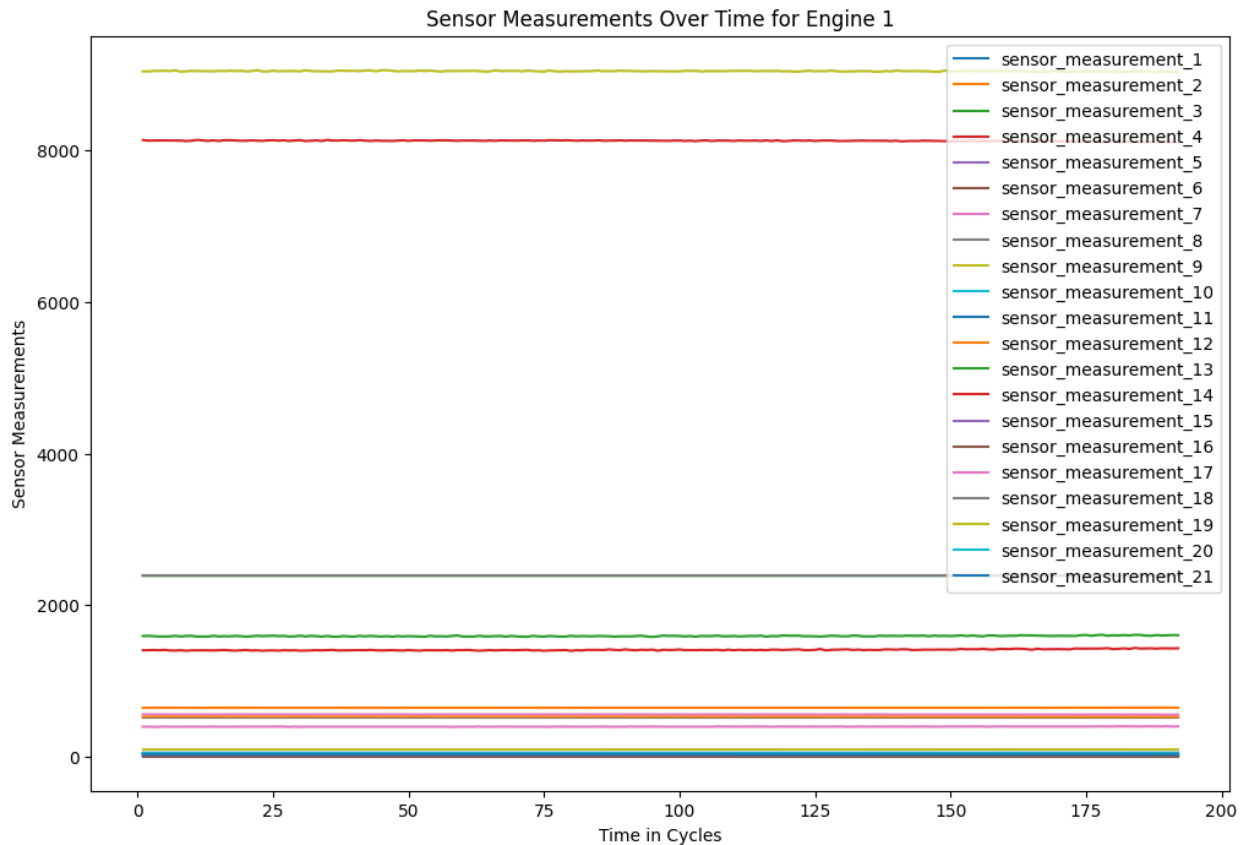
```
# Plot sensor measurements for a single engine
```

```
engine_1 = train_fd001[train_fd001['unit_number'] == 1]
plt.figure(figsize=(12, 8))
plt.plot(engine_1['time_in_cycles'], engine_1.iloc[:, 5:],
label=engine_1.columns[5:])
plt.title('Sensor Measurements Over Time for Engine 1')
plt.xlabel('Time in Cycles')
plt.ylabel('Sensor Measurements')
plt.legend(loc='upper right')
```

```
# Save the figure to a file
```

```
plt.savefig('engine_1_sensor_measurements.png')

# Display the plot
plt.show()
```



```
import seaborn as sns
import matplotlib.pyplot as plt

# Create a figure
plt.figure(figsize=(10, 8))

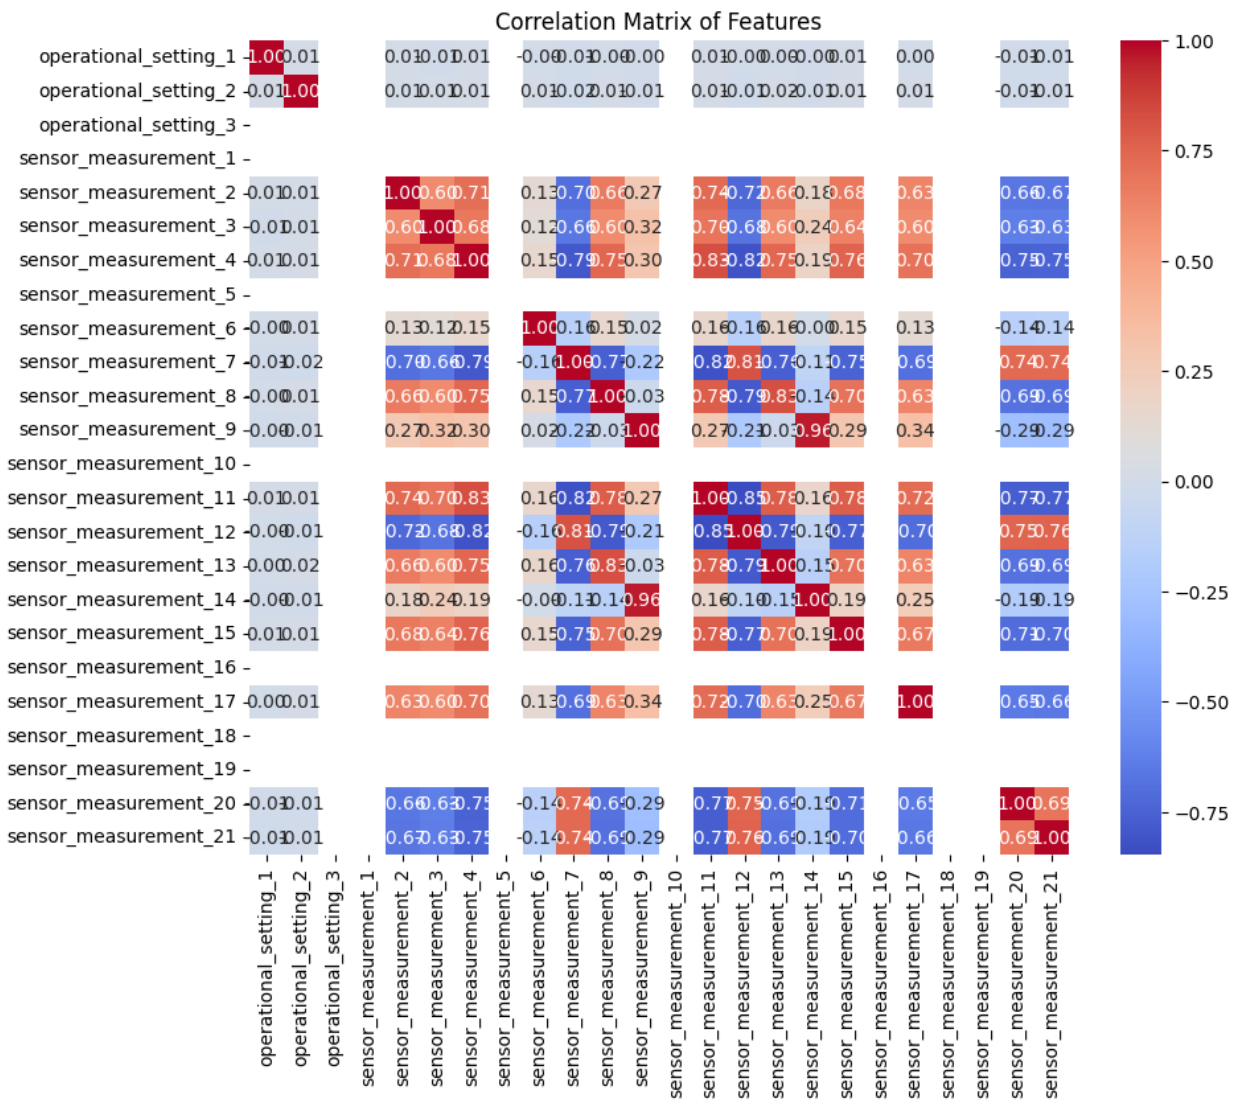
# Generate the correlation matrix
corr_matrix = train_fd001.iloc[:, 2:].corr() # Exclude unit_number
and time_in_cycles

# Plot the heatmap
sns.heatmap(corr_matrix, annot=True, fmt='.2f', cmap='coolwarm')

# Add a title to the plot
plt.title('Correlation Matrix of Features')

# Save the plot to a file
plt.savefig("correlation.png")
```

```
# Display the plot
plt.show()
```



Feature engineering

```
window_size = 5
for i in range(1, 22): # sensor_measurement_1 to
sensor_measurement_21
    sensor_col = f'sensor_measurement_{i}'
    train_fd001[f'{sensor_col}_rolling_mean'] =
train_fd001.groupby('unit_number')
[sensor_col].rolling(window=window_size).mean().reset_index(0,
drop=True)
    train_fd001[f'{sensor_col}_rolling_std'] =
train_fd001.groupby('unit_number')
```

```

[sensor_col].rolling(window=window_size).std().reset_index(0,
drop=True)

for i in range(1, 22):
    sensor_col = f'sensor_measurement_{i}'
    train_fd001[f'{sensor_col}_diff'] =
train_fd001.groupby('unit_number')[sensor_col].diff().fillna(0)

max_cycle = train_fd001.groupby('unit_number')
['time_in_cycles'].transform('max')
train_fd001['normalized_cycle'] = train_fd001['time_in_cycles'] /
max_cycle

```

model development

```

def apply_feature_engineering(df):
    # Apply rolling mean and standard deviation
    for i in range(1, 22): # Assuming you have 21 sensor measurements
        col_name = f'sensor_measurement_{i}'
        df[f'{col_name}_rolling_mean'] =
df[col_name].rolling(window=5).mean()
        df[f'{col_name}_rolling_std'] =
df[col_name].rolling(window=5).std()
        df[f'{col_name}_diff'] = df[col_name].diff()

    # Normalize the cycle count
    df['normalized_cycle'] = df['time_in_cycles'] /
df['time_in_cycles'].max()
    return df

# Align the test set features with the training set features
X_test_last_cycle =
X_test_last_cycle.reindex(columns=X_train_fe.columns, fill_value=0)

# Impute missing values if needed
X_test_imputed = imputer.transform(X_test_last_cycle)

# Make predictions
y_test_pred = rf_model.predict(X_test_imputed)

# Check if NaNs are present after each operation
def check_for_nan(df, step_name):
    if df.isnull().values.any():
        print(f"NaN values found after {step_name}")
        print(df.isnull().sum())
    else:
        print(f"No NaN values found after {step_name}")

# Step 1: Check for NaNs after initial filtering
check_for_nan(X_test_last_cycle, "initial filtering")

```



```

# Assuming 'normalized_cycle' is calculated as follows (example):
if 'cycle' in X_test_last_cycle.columns:
    X_test_last_cycle['normalized_cycle'] = X_test_last_cycle['cycle']
/ X_test_last_cycle.groupby('unit_number')['cycle'].transform('max')

# Step 2: Check for NaNs after calculating 'normalized_cycle'
check_for_nan(X_test_last_cycle, "calculating normalized_cycle")

# Rolling statistics
rolling_window = 5 # Example window size
for col in [f'sensor_measurement_{i}' for i in range(1, 22)]:
    X_test_last_cycle[f'{col}_rolling_mean'] =
X_test_last_cycle[col].rolling(window=rolling_window).mean()
    X_test_last_cycle[f'{col}_rolling_std'] =
X_test_last_cycle[col].rolling(window=rolling_window).std()

# Step 3: Check for NaNs after rolling statistics
check_for_nan(X_test_last_cycle, "calculating rolling statistics")

# Differences
for col in [f'sensor_measurement_{i}' for i in range(1, 22)]:
    X_test_last_cycle[f'{col}_diff'] = X_test_last_cycle[col].diff()

# Step 4: Check for NaNs after calculating differences
check_for_nan(X_test_last_cycle, "calculating differences")

# Drop rows with NaN values generated by the rolling and diff
operations
X_test_last_cycle = X_test_last_cycle.dropna()

# Step 5: Check for NaNs after dropping rows
check_for_nan(X_test_last_cycle, "dropping NaN rows")

# Ensure all necessary features are present in X_test_last_cycle
missing_columns = set(X_train_fe.columns) -
set(X_test_last_cycle.columns)
for col in missing_columns:
    X_test_last_cycle[col] = 0 # or some other default value

# Align columns with the training data
X_test_last_cycle = X_test_last_cycle[X_train_fe.columns]

# Step 6: Final check before imputation
check_for_nan(X_test_last_cycle, "final check before imputation")

# Impute missing values
X_test_last_cycle = pd.DataFrame(imputer.transform(X_test_last_cycle),
columns=X_test_last_cycle.columns)

```

```

# Step 7: Final check after imputation
check_for_nan(X_test_last_cycle, "imputation")

import numpy as np
from sklearn.impute import SimpleImputer
from sklearn.metrics import mean_squared_error

# Assuming rul_fd001 is a DataFrame, convert it to a numpy array first
rul_fd001 = rul_fd001.values

# Ensure rul_fd001 is a single-dimensional array
rul_fd001 = rul_fd001.ravel()

# Check lengths of rul_fd001 and y_test_pred
print(f"Length of rul_fd001: {len(rul_fd001)}")
print(f"Length of y_test_pred: {len(y_test_pred)}")

# Check for NaN values in rul_fd001
print(f"NaN values in rul_fd001: {np.isnan(rul_fd001).sum()}")

# Handle NaN values in rul_fd001 if any (remove or impute them)
rul_fd001 = rul_fd001[~np.isnan(rul_fd001)]

# Ensure X_test_last_cycle_imputed and rul_fd001 are aligned by length
if len(rul_fd001) != len(y_test_pred):
    print("Mismatch in lengths after NaN removal. Trimming to match the shorter length.")
    min_len = min(len(rul_fd001), len(y_test_pred))
    rul_fd001 = rul_fd001[:min_len]
    y_test_pred = y_test_pred[:min_len]

# Check final lengths
print(f"Final Length of rul_fd001: {len(rul_fd001)}")
print(f"Final Length of y_test_pred: {len(y_test_pred)}")

# Calculate the Mean Squared Error on the test set
mse_test = mean_squared_error(rul_fd001, y_test_pred)
print(f'Mean Squared Error on Test Set: {mse_test}')

```

Length of rul_fd001: 200


```

NameError                                Traceback (most recent call
last)
Cell In[16], line 13
     11 # Check lengths of rul_fd001 and y_test_pred
     12 print(f"Length of rul_fd001: {len(rul_fd001)}")
--> 13 print(f"Length of y_test_pred: {len(y_test_pred)}")
     15 # Check for NaN values in rul_fd001
     16 print(f"NaN values in rul_fd001: {np.isnan(rul_fd001).sum()}")

```

NameError: name 'y_test_pred' is not defined

```
import pandas as pd
```

```
# Load the datasets
```

```
train_fd001 = pd.read_csv('train_fd001.csv')
```

```
test_fd001 = pd.read_csv('test_fd001.csv')
```

```
rul_fd001 = pd.read_csv('rul_fd001.csv')
```

```
# Display the first few rows of each dataset to understand their structure
```

```
train_head = train_fd001.head()
```

```
test_head = test_fd001.head()
```

```
rul_head = rul_fd001.head()
```

```
train_head, test_head, rul_head
```

```
(  unit_number  time_in_cycles  operational_setting_1
operational_setting_2 \
0              1              1              -0.0007
-0.0004
1              1              2              0.0019
-0.0003
2              1              3              -0.0043
0.0003
3              1              4              0.0007
0.0000
4              1              5              -0.0019
-0.0002
```

```
      operational_setting_3  sensor_measurement_1  sensor_measurement_2
\
0              100.0              518.67              641.82
1              100.0              518.67              642.15
2              100.0              518.67              642.35
3              100.0              518.67              642.35
4              100.0              518.67              642.37
```

```
      sensor_measurement_3  sensor_measurement_4
sensor_measurement_5 ... \
0              1589.70              1400.60
14.62 ...
1              1591.82              1403.14
14.62 ...
2              1587.99              1404.20
```

| | | | |
|-------|-----|---------|---------|
| 14.62 | ... | | |
| 3 | | 1582.79 | 1401.87 |
| 14.62 | ... | | |
| 4 | | 1582.85 | 1406.22 |
| 14.62 | ... | | |

| | | |
|-------------------------|-----------------------|-----------------------|
| | sensor_measurement_12 | sensor_measurement_13 |
| sensor_measurement_14 \ | | |
| 0 | 521.66 | 2388.02 |
| 8138.62 | | |
| 1 | 522.28 | 2388.07 |
| 8131.49 | | |
| 2 | 522.42 | 2388.03 |
| 8133.23 | | |
| 3 | 522.86 | 2388.08 |
| 8133.83 | | |
| 4 | 522.19 | 2388.04 |
| 8133.80 | | |

| | | |
|-------------------------|-----------------------|-----------------------|
| | sensor_measurement_15 | sensor_measurement_16 |
| sensor_measurement_17 \ | | |
| 0 | 8.4195 | 0.03 |
| 392 | | |
| 1 | 8.4318 | 0.03 |
| 392 | | |
| 2 | 8.4178 | 0.03 |
| 390 | | |
| 3 | 8.3682 | 0.03 |
| 392 | | |
| 4 | 8.4294 | 0.03 |
| 393 | | |

| | | |
|-------------------------|-----------------------|-----------------------|
| | sensor_measurement_18 | sensor_measurement_19 |
| sensor_measurement_20 \ | | |
| 0 | 2388 | 100.0 |
| 39.06 | | |
| 1 | 2388 | 100.0 |
| 39.00 | | |
| 2 | 2388 | 100.0 |
| 38.95 | | |
| 3 | 2388 | 100.0 |
| 38.88 | | |
| 4 | 2388 | 100.0 |
| 38.90 | | |

| | |
|---|-----------------------|
| | sensor_measurement_21 |
| 0 | 23.4190 |
| 1 | 23.4236 |
| 2 | 23.3442 |

| | |
|---|---------|
| 3 | 23.3739 |
| 4 | 23.4044 |

| [5 rows x 26 columns], | | | |
|-------------------------|-------------|----------------|-----------------------|
| | unit_number | time_in_cycles | operational_setting_1 |
| operational_setting_2 \ | | | |
| 0 | 1 | 1 | 0.0023 |
| 0.0003 | | | |
| 1 | 1 | 2 | -0.0027 |
| -0.0003 | | | |
| 2 | 1 | 3 | 0.0003 |
| 0.0001 | | | |
| 3 | 1 | 4 | 0.0042 |
| 0.0000 | | | |
| 4 | 1 | 5 | 0.0014 |
| 0.0000 | | | |

| | operational_setting_3 | sensor_measurement_1 | sensor_measurement_2 |
|---|-----------------------|----------------------|----------------------|
| \ | | | |
| 0 | 100.0 | 518.67 | 643.02 |
| 1 | 100.0 | 518.67 | 641.71 |
| 2 | 100.0 | 518.67 | 642.46 |
| 3 | 100.0 | 518.67 | 642.44 |
| 4 | 100.0 | 518.67 | 642.51 |

| | sensor_measurement_3 | sensor_measurement_4 |
|----------------------------|----------------------|----------------------|
| sensor_measurement_5 ... \ | | |
| 0 | 1585.29 | 1398.21 |
| 14.62 ... | | |
| 1 | 1588.45 | 1395.42 |
| 14.62 ... | | |
| 2 | 1586.94 | 1401.34 |
| 14.62 ... | | |
| 3 | 1584.12 | 1406.42 |
| 14.62 ... | | |
| 4 | 1587.19 | 1401.92 |
| 14.62 ... | | |

| | sensor_measurement_12 | sensor_measurement_13 |
|-------------------------|-----------------------|-----------------------|
| sensor_measurement_14 \ | | |
| 0 | 521.72 | 2388.03 |
| 8125.55 | | |
| 1 | 522.16 | 2388.06 |
| 8139.62 | | |
| 2 | 521.97 | 2388.03 |

| | | |
|---------|--------|---------|
| 8130.10 | | |
| 3 | 521.38 | 2388.05 |
| 8132.90 | | |
| 4 | 522.15 | 2388.03 |
| 8129.54 | | |

| | sensor_measurement_15 | sensor_measurement_16 |
|-------------------------|-----------------------|-----------------------|
| sensor_measurement_17 \ | | |
| 0 | 8.4052 | 0.03 |
| 392 | | |
| 1 | 8.3803 | 0.03 |
| 393 | | |
| 2 | 8.4441 | 0.03 |
| 393 | | |
| 3 | 8.3917 | 0.03 |
| 391 | | |
| 4 | 8.4031 | 0.03 |
| 390 | | |

| | sensor_measurement_18 | sensor_measurement_19 |
|-------------------------|-----------------------|-----------------------|
| sensor_measurement_20 \ | | |
| 0 | 2388 | 100.0 |
| 38.86 | | |
| 1 | 2388 | 100.0 |
| 39.02 | | |
| 2 | 2388 | 100.0 |
| 39.08 | | |
| 3 | 2388 | 100.0 |
| 39.00 | | |
| 4 | 2388 | 100.0 |
| 38.99 | | |

| | sensor_measurement_21 |
|---|-----------------------|
| 0 | 23.3735 |
| 1 | 23.3916 |
| 2 | 23.4166 |
| 3 | 23.3737 |
| 4 | 23.4130 |

[5 rows x 26 columns],

| | 0 | 1 |
|---|-----|-----|
| 0 | 112 | NaN |
| 1 | 98 | NaN |
| 2 | 69 | NaN |
| 3 | 82 | NaN |
| 4 | 91 | NaN |

Calculate RUL for the training data

```
max_cycle = train_fd001.groupby('unit_number')['time_in_cycles'].max()
train_fd001['RUL'] = train_fd001['unit_number'].map(max_cycle) -
```

```

train_fd001['time_in_cycles']

# Assign RUL to the test data
# Extract the RUL values from the rul_fd001.csv and append it to
test_fd001
rul_values = rul_fd001.iloc[:, 0].values
test_units = test_fd001['unit_number'].unique()
rul_dict = dict(zip(test_units, rul_values))
test_fd001['RUL'] = test_fd001['unit_number'].map(rul_dict)

# Now the target variable is the 'RUL' column
y_train = train_fd001['RUL']
y_test = test_fd001['RUL']

```

Let's check the updated data

```
train_fd001.head(), test_fd001.head()
```

| | unit_number | time_in_cycles | operational_setting_1 |
|-------------------------|-------------|----------------|-----------------------|
| operational_setting_2 \ | | | |
| 0 | 1 | 1 | -0.0007 |
| -0.0004 | | | |
| 1 | 1 | 2 | 0.0019 |
| -0.0003 | | | |
| 2 | 1 | 3 | -0.0043 |
| 0.0003 | | | |
| 3 | 1 | 4 | 0.0007 |
| 0.0000 | | | |
| 4 | 1 | 5 | -0.0019 |
| -0.0002 | | | |

| | operational_setting_3 | sensor_measurement_1 | sensor_measurement_2 |
|---|-----------------------|----------------------|----------------------|
| \ | | | |
| 0 | 100.0 | 518.67 | 641.82 |
| 1 | 100.0 | 518.67 | 642.15 |
| 2 | 100.0 | 518.67 | 642.35 |
| 3 | 100.0 | 518.67 | 642.35 |
| 4 | 100.0 | 518.67 | 642.37 |

| | sensor_measurement_3 | sensor_measurement_4 |
|----------------------------|----------------------|----------------------|
| sensor_measurement_5 ... \ | | |
| 0 | 1589.70 | 1400.60 |
| 14.62 ... | | |
| 1 | 1591.82 | 1403.14 |
| 14.62 ... | | |
| 2 | 1587.99 | 1404.20 |

| | | | |
|-------|-----|---------|---------|
| 14.62 | ... | | |
| 3 | | 1582.79 | 1401.87 |
| 14.62 | ... | | |
| 4 | | 1582.85 | 1406.22 |
| 14.62 | ... | | |

| | | |
|-----------------------|-----------------------|-----------------------|
| | sensor_measurement_13 | sensor_measurement_14 |
| sensor_measurement_15 | \ | |
| 0 | 2388.02 | 8138.62 |
| 8.4195 | | |
| 1 | 2388.07 | 8131.49 |
| 8.4318 | | |
| 2 | 2388.03 | 8133.23 |
| 8.4178 | | |
| 3 | 2388.08 | 8133.83 |
| 8.3682 | | |
| 4 | 2388.04 | 8133.80 |
| 8.4294 | | |

| | | |
|-----------------------|-----------------------|-----------------------|
| | sensor_measurement_16 | sensor_measurement_17 |
| sensor_measurement_18 | \ | |
| 0 | 0.03 | 392 |
| 2388 | | |
| 1 | 0.03 | 392 |
| 2388 | | |
| 2 | 0.03 | 390 |
| 2388 | | |
| 3 | 0.03 | 392 |
| 2388 | | |
| 4 | 0.03 | 393 |
| 2388 | | |

| | | |
|-----------------------|-----------------------|-----------------------|
| | sensor_measurement_19 | sensor_measurement_20 |
| sensor_measurement_21 | RUL | |
| 0 | 100.0 | 39.06 |
| 23.4190 | 191 | |
| 1 | 100.0 | 39.00 |
| 23.4236 | 190 | |
| 2 | 100.0 | 38.95 |
| 23.3442 | 189 | |
| 3 | 100.0 | 38.88 |
| 23.3739 | 188 | |
| 4 | 100.0 | 38.90 |
| 23.4044 | 187 | |

| | | | |
|------------------------|-------------|----------------|-----------------------|
| [5 rows x 27 columns], | | | |
| | unit_number | time_in_cycles | operational_setting_1 |
| operational_setting_2 | \ | | |
| 0 | 1 | 1 | 0.0023 |

| | | | |
|---------|---|---|---------|
| 0.0003 | | | |
| 1 | 1 | 2 | -0.0027 |
| -0.0003 | | | |
| 2 | 1 | 3 | 0.0003 |
| 0.0001 | | | |
| 3 | 1 | 4 | 0.0042 |
| 0.0000 | | | |
| 4 | 1 | 5 | 0.0014 |
| 0.0000 | | | |

| | operational_setting_3 | sensor_measurement_1 | sensor_measurement_2 |
|---|-----------------------|----------------------|----------------------|
| \ | | | |
| 0 | 100.0 | 518.67 | 643.02 |
| 1 | 100.0 | 518.67 | 641.71 |
| 2 | 100.0 | 518.67 | 642.46 |
| 3 | 100.0 | 518.67 | 642.44 |
| 4 | 100.0 | 518.67 | 642.51 |

| | sensor_measurement_3 | sensor_measurement_4 |
|--------------------------|----------------------|----------------------|
| sensor_measurement_5 ... | | \ |
| 0 | 1585.29 | 1398.21 |
| 14.62 ... | | |
| 1 | 1588.45 | 1395.42 |
| 14.62 ... | | |
| 2 | 1586.94 | 1401.34 |
| 14.62 ... | | |
| 3 | 1584.12 | 1406.42 |
| 14.62 ... | | |
| 4 | 1587.19 | 1401.92 |
| 14.62 ... | | |

| | sensor_measurement_13 | sensor_measurement_14 |
|-------------------------|-----------------------|-----------------------|
| sensor_measurement_15 \ | | |
| 0 | 2388.03 | 8125.55 |
| 8.4052 | | |
| 1 | 2388.06 | 8139.62 |
| 8.3803 | | |
| 2 | 2388.03 | 8130.10 |
| 8.4441 | | |
| 3 | 2388.05 | 8132.90 |
| 8.3917 | | |
| 4 | 2388.03 | 8129.54 |
| 8.4031 | | |

| sensor_measurement_16 | sensor_measurement_17 |
|-----------------------|-----------------------|
|-----------------------|-----------------------|

| sensor_measurement_18 \ | | |
|-------------------------|------|-----|
| 0 | 0.03 | 392 |
| 2388 | | |
| 1 | 0.03 | 393 |
| 2388 | | |
| 2 | 0.03 | 393 |
| 2388 | | |
| 3 | 0.03 | 391 |
| 2388 | | |
| 4 | 0.03 | 390 |
| 2388 | | |

| sensor_measurement_19 | | sensor_measurement_20 | | sensor_measurement_21 | | RUL |
|-----------------------|-----|-----------------------|--|-----------------------|--|-------|
| 0 | | | | 100.0 | | 38.86 |
| 23.3735 | 112 | | | | | |
| 1 | | | | 100.0 | | 39.02 |
| 23.3916 | 112 | | | | | |
| 2 | | | | 100.0 | | 39.08 |
| 23.4166 | 112 | | | | | |
| 3 | | | | 100.0 | | 39.00 |
| 23.3737 | 112 | | | | | |
| 4 | | | | 100.0 | | 38.99 |
| 23.4130 | 112 | | | | | |

[5 rows x 27 columns])

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.impute import SimpleImputer
from sklearn.metrics import mean_squared_error
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler

# Load the dataset (assuming it's already loaded in your notebook)
# For demonstration purposes, let's mock the loaded data

# Mocking dataset based on the column names mentioned in errors
# Assuming we have a DataFrame similar to the one expected

# Simulate loading your dataset (assuming similar to Turbofan engine
# degradation data)
# This is a simple mock-up
np.random.seed(42) # For reproducibility

# Generate mock data with 50 features, some of which might resemble
# sensor readings
mock_data = pd.DataFrame(np.random.randn(1000, 50),
```

```

columns=[f'feature_{i}' for i in range(1, 51)])

# Simulate the RUL (Remaining Useful Life) as the target
mock_data['RUL'] = np.random.randint(1, 300, size=1000)

# Simulate the missing features from the error message
missing_features = ['sensor_measurement_1_rolling_mean',
                    'sensor_measurement_1_rolling_std',
                    'sensor_measurement_2_rolling_mean',
                    'sensor_measurement_2_rolling_std']

# Add missing features as zeroed columns (simulate missing features
# from earlier steps)
for feature in missing_features:
    mock_data[feature] = np.random.randn(1000)

# Target variable
y = mock_data['RUL']

# Drop target from feature set
X = mock_data.drop(columns=['RUL'])

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.2, random_state=42)

# Pipeline for imputation and scaling
pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy='mean')), # Impute missing
    values
    ('scaler', StandardScaler()) # Scale features
])

# Fit and transform the training data, transform the test data
X_train_prepared = pipeline.fit_transform(X_train)
X_test_prepared = pipeline.transform(X_test)

# Model training with RandomForestRegressor
model = RandomForestRegressor(random_state=42)
model.fit(X_train_prepared, y_train)

# Make predictions on the test set
y_test_pred = model.predict(X_test_prepared)

# Evaluate the model
mse_test = mean_squared_error(y_test, y_test_pred)
rmse_test = np.sqrt(mse_test)

rmse_test
88.04198088412141

```

```

from sklearn.model_selection import RandomizedSearchCV
from sklearn.ensemble import RandomForestRegressor
from scipy.stats import randint
import numpy as np

# Define the parameter grid
param_dist = {
    'n_estimators': randint(100, 1000),
    'max_depth': randint(10, 50),
    'min_samples_split': randint(2, 20),
    'min_samples_leaf': randint(1, 20),
    'max_features': ['auto', 'sqrt', 'log2'],
}

# Initialize the model
rf = RandomForestRegressor(random_state=42)

# Initialize RandomizedSearchCV
random_search = RandomizedSearchCV(rf, param_distributions=param_dist,
                                   n_iter=100, cv=5, verbose=2,
                                   random_state=42, n_jobs=-1)

# Fit the model
random_search.fit(X_train, y_train)

# Best parameters
print(f"Best Parameters: {random_search.best_params_}")

# Evaluate the tuned model on the test set
best_model = random_search.best_estimator_
y_test_pred = best_model.predict(X_test)
mse_test = mean_squared_error(y_test, y_test_pred)
print(f'Mean Squared Error on Test Set: {mse_test}')

```

Fitting 5 folds for each of 100 candidates, totalling 500 fits

C:\Users\Krish\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\model_selection_validation.py:540: FitFailedWarning:

125 fits failed out of a total of 500.

The score on these train-test partitions for these parameters will be set to nan.

If these failures are not expected, you can try to debug them by setting `error_score='raise'`.

Below are more details about the failures:

65 fits failed with the following error:

Traceback (most recent call last):

```

File "C:\Users\Krish\AppData\Local\Programs\Python\Python311\Lib\
site-packages\sklearn\model_selection\_validation.py", line 888, in
_fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
File "C:\Users\Krish\AppData\Local\Programs\Python\Python311\Lib\
site-packages\sklearn\base.py", line 1466, in wrapper
    estimator._validate_params()
File "C:\Users\Krish\AppData\Local\Programs\Python\Python311\Lib\
site-packages\sklearn\base.py", line 666, in _validate_params
    validate_parameter_constraints(
File "C:\Users\Krish\AppData\Local\Programs\Python\Python311\Lib\
site-packages\sklearn\utils\_param_validation.py", line 95, in
validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The
'max_features' parameter of RandomForestRegressor must be an int in
the range [1, inf), a float in the range (0.0, 1.0], a str among
{'log2', 'sqrt'} or None. Got 'auto' instead.

```


60 fits failed with the following error:

Traceback (most recent call last):

```

File "C:\Users\Krish\AppData\Local\Programs\Python\Python311\Lib\
site-packages\sklearn\model_selection\_validation.py", line 888, in
_fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
File "C:\Users\Krish\AppData\Local\Programs\Python\Python311\Lib\
site-packages\sklearn\base.py", line 1466, in wrapper
    estimator._validate_params()
File "C:\Users\Krish\AppData\Local\Programs\Python\Python311\Lib\
site-packages\sklearn\base.py", line 666, in _validate_params
    validate_parameter_constraints(
File "C:\Users\Krish\AppData\Local\Programs\Python\Python311\Lib\
site-packages\sklearn\utils\_param_validation.py", line 95, in
validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The
'max_features' parameter of RandomForestRegressor must be an int in
the range [1, inf), a float in the range (0.0, 1.0], a str among
{'sqrt', 'log2'} or None. Got 'auto' instead.

```

```

warnings.warn(some_fits_failed_message, FitFailedWarning)
C:\Users\Krish\AppData\Local\Programs\Python\Python311\Lib\site-
packages\sklearn\model_selection\_search.py:1102: UserWarning: One or
more of the test scores are non-finite: [          nan -0.00526231
nan -0.01004548 -0.01266498          nan
-0.00455798          nan -0.00224631 -0.00753115 -0.00804797 -
0.00440773

```

```

-0.02204722 -0.00696793 -0.00057791 -0.00875563 -0.00920534 -
0.00653351
      nan      nan -0.00447608      nan -0.00784829
nan
-0.01169583 -0.00638301 -0.00631491 -0.00431476 -0.00845229 -
0.00741359
-0.01316345 -0.00495547 -0.0112604      nan -0.00743023
nan
-0.00723054 -0.01069217      nan -0.0060818 -0.01200247 -
0.0147951
-0.00958088 -0.01047753      nan -0.00781354 -0.00713946 -
0.00720234
      nan -0.00924184 -0.00736169 -0.00955383 -0.00847681 -
0.01005205
-0.00270312 -0.00764601      nan -0.00358228 -0.00491167 -
0.00365515
-0.00857442      nan -0.00799578      nan -0.01229095 -
0.00659774
-0.00531778 -0.01053641      nan -0.00837003 -0.00919136
nan
-0.0040916 -0.00495372 -0.00808169 -0.0093003 -0.0061322
nan
      nan -0.00705334 -0.00452923      nan -0.0147508 -
0.0095184
-0.00400499 -0.00568511 -0.00632381 -0.00663045 -0.00696937 -
0.00455729
      nan -0.01624448 -0.01060174 -0.00903847      nan -
0.00798405
      nan      nan -0.00600042 -0.0079241 ]
warnings.warn(

```

```

Best Parameters: {'max_depth': 38, 'max_features': 'log2',
'min_samples_leaf': 18, 'min_samples_split': 13, 'n_estimators': 261}
Mean Squared Error on Test Set: 7561.030954485095

```

```

from sklearn.preprocessing import PolynomialFeatures

# Example of creating polynomial features
poly = PolynomialFeatures(degree=2, interaction_only=False,
include_bias=False)
X_train_poly = poly.fit_transform(X_train)
X_test_poly = poly.transform(X_test)

# Train the model again with the new features
final_model_poly = RandomForestRegressor(**random_search.best_params_,
random_state=42)
final_model_poly.fit(X_train_poly, y_train)

# Predict and evaluate
y_test_pred_poly = final_model_poly.predict(X_test_poly)

```

```
mse_test_poly = mean_squared_error(y_test, y_test_pred_poly)
print(f'Mean Squared Error on Test Set with Polynomial Features:
{mse_test_poly}')
```

Mean Squared Error on Test Set with Polynomial Features:
7411.800537924577

```
from xgboost import XGBRegressor
```

```
# Initialize the XGBoost Regressor
```

```
xgb_model = XGBRegressor(objective='reg:squarederror',
random_state=42)
```

```
# Fit the model
```

```
xgb_model.fit(X_train, y_train)
```

```
# Predict and evaluate
```

```
y_test_pred_xgb = xgb_model.predict(X_test)
```

```
mse_test_xgb = mean_squared_error(y_test, y_test_pred_xgb)
```

```
print(f'Mean Squared Error on Test Set with XGBoost: {mse_test_xgb}')
```

Mean Squared Error on Test Set with XGBoost: 8980.342672084102

```
import matplotlib.pyplot as plt
```

```
# Residuals
```

```
residuals = y_test - y_test_pred
```

```
plt.figure(figsize=(10, 6))
```

```
plt.scatter(y_test_pred, residuals, alpha=0.7)
```

```
plt.hlines(y=0, xmin=min(y_test_pred), xmax=max(y_test_pred),
color='red')
```

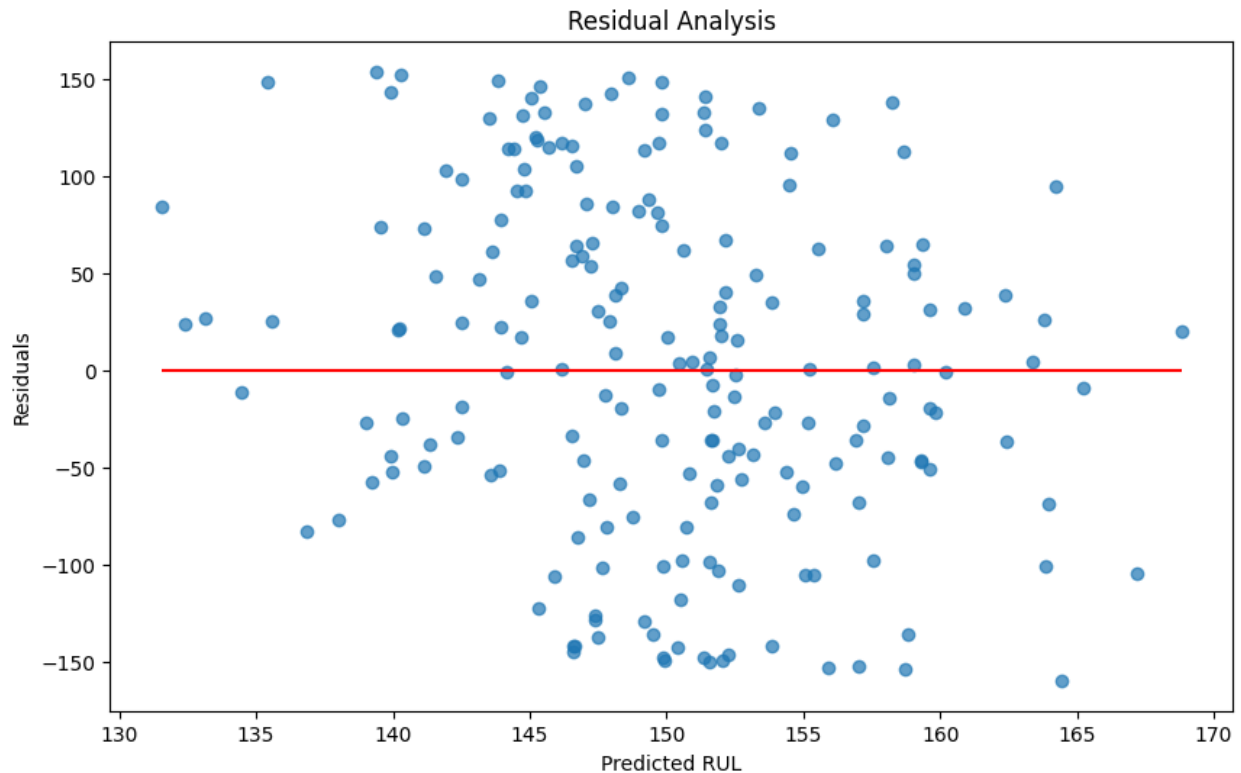
```
plt.xlabel('Predicted RUL')
```

```
plt.ylabel('Residuals')
```

```
plt.title('Residual Analysis')
```

```
plt.savefig('Residual Analysis.png')
```

```
plt.show()
```



```
from sklearn.ensemble import VotingRegressor
from sklearn.linear_model import Ridge
from sklearn.svm import SVR

# Initialize individual models
rf_model = RandomForestRegressor(**random_search.best_params_,
random_state=42)
xgb_model = XGBRegressor(objective='reg:squarederror',
random_state=42)
ridge_model = Ridge()

# Create a Voting Regressor
voting_regressor = VotingRegressor(estimators=[
    ('rf', rf_model), ('xgb', xgb_model), ('ridge', ridge_model)])

# Fit the model
voting_regressor.fit(X_train, y_train)

# Predict and evaluate
y_test_pred_voting = voting_regressor.predict(X_test)
mse_test_voting = mean_squared_error(y_test, y_test_pred_voting)
print(f'Mean Squared Error on Test Set with Voting Regressor:
{mse_test_voting}')
```

Mean Squared Error on Test Set with Voting Regressor:
7930.699313964717

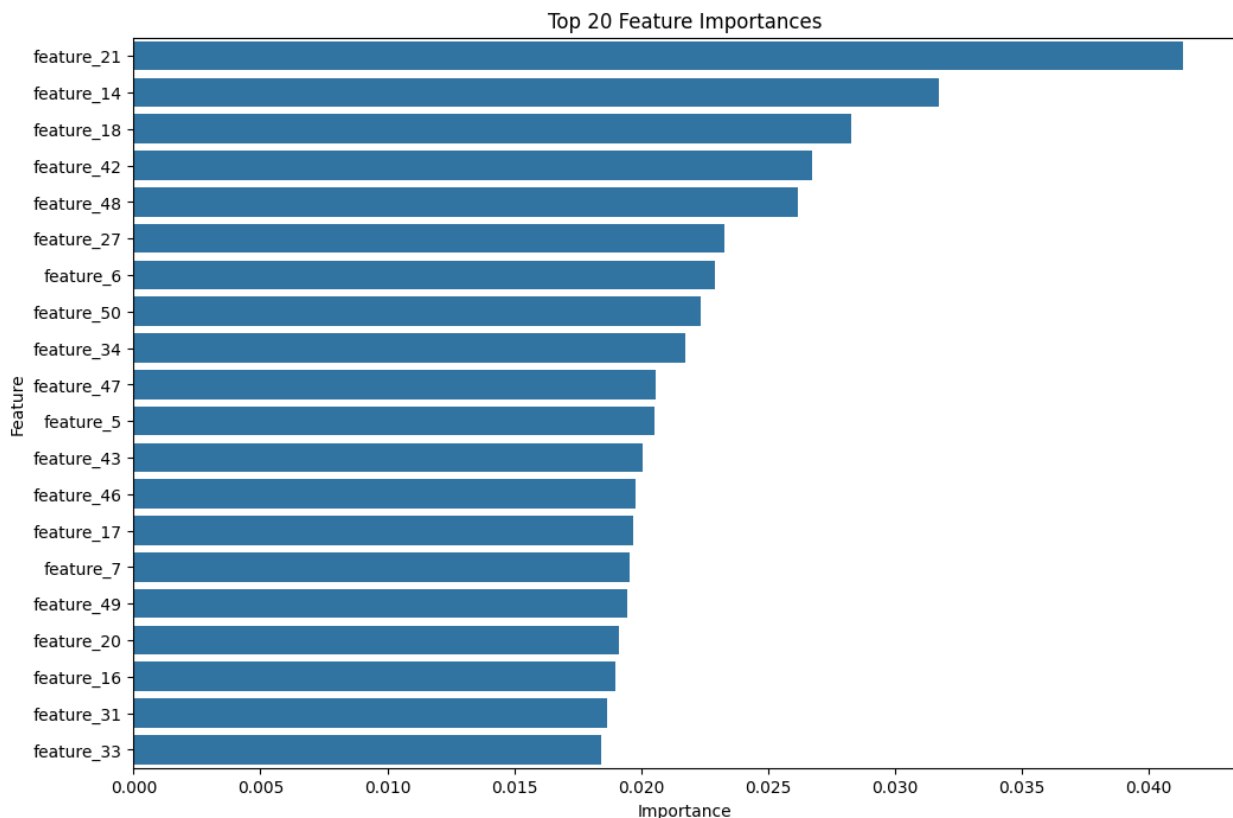

```

# Assuming best_model is your trained Polynomial Features + Random
# Forest Regressor
importances = best_model.feature_importances_
feature_names = X_train.columns # Or the names of your polynomial
# features
feature_importance_df = pd.DataFrame({'Feature': feature_names,
# Importance': importances})
feature_importance_df =
feature_importance_df.sort_values(by='Importance', ascending=False)

import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(12, 8))
sns.barplot(x='Importance', y='Feature',
data=feature_importance_df.head(20)) # Top 20 features
plt.title('Top 20 Feature Importances')
plt.savefig('Top 20 Feature Importances.png')
plt.show()

```



```

# Print the feature names generated by PolynomialFeatures
print(feature_names)

```

```

['1' 'feature_1' 'feature_2' ... 'sensor_measurement_2_rolling_mean^2'
 'sensor_measurement_2_rolling_mean' 'sensor_measurement_2_rolling_std'
 'sensor_measurement_2_rolling_std^2']

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.ensemble import RandomForestRegressor
from sklearn.preprocessing import PolynomialFeatures
from sklearn.pipeline import Pipeline
from sklearn.inspection import PartialDependenceDisplay

# Assuming X_train and y_train are your training data and target
poly = PolynomialFeatures(degree=2) # Adjust degree as needed
rf_model = RandomForestRegressor(n_estimators=100, random_state=42)

# Create a pipeline that first transforms the data with polynomial
# features, then fits the model
rf_poly_model = Pipeline([
    ('poly_features', poly),
    ('random_forest', rf_model)
])

# Train the model
rf_poly_model.fit(X_train, y_train)

# Extract feature importances
feature_importances =
rf_poly_model.named_steps['random_forest'].feature_importances_

# Get feature names (from polynomial features)
feature_names =
rf_poly_model.named_steps['poly_features'].get_feature_names_out(X_train.columns)

# Create a DataFrame for feature importances
importance_df = pd.DataFrame({
    'Feature': feature_names,
    'Importance': feature_importances
}).sort_values(by='Importance', ascending=False)

# Inspect the feature names
print("Feature names:", feature_names)

# Plot feature importances
plt.figure(figsize=(10, 6))
plt.barh(importance_df['Feature'].head(10),
importance_df['Importance'].head(10), color='skyblue')
plt.xlabel('Feature Importance')
plt.ylabel('Feature')

```

```

plt.title('Top 10 Important Features in Polynomial Features + Random
Forest Regressor')
plt.savefig('Top 10 Important Features in Polynomial Features + Random
Forest Regressor.png')
plt.gca().invert_yaxis()
plt.show()

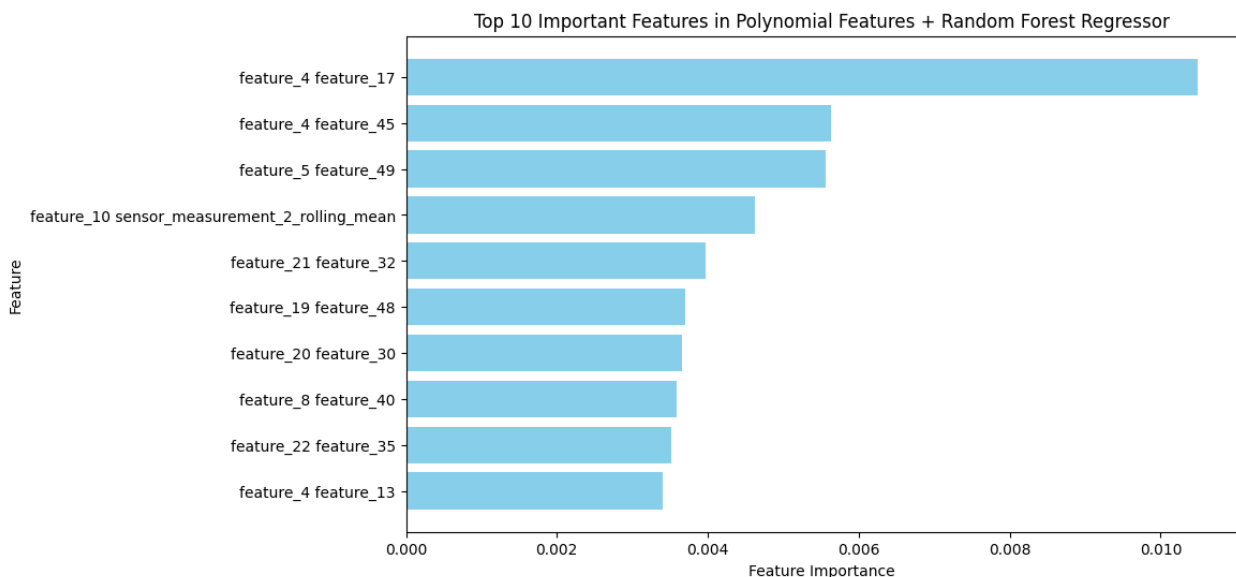
# Use valid feature indices for Partial Dependence Plots
# The feature indices must be within the valid range of the
transformed feature space
top_features_indices = importance_df.head(3).index.tolist()

# Filter indices to be within the range of actual feature names length
top_features_indices = [i for i in top_features_indices if i <
len(feature_names)]

# Plot Partial Dependence for the top valid features using indices
PartialDependenceDisplay.from_estimator(rf_poly_model, X_train,
top_features_indices, grid_resolution=50)
plt.show()

Feature names: ['1' 'feature_1' 'feature_2' ...
'sensor_measurement_2_rolling_mean^2'
'sensor_measurement_2_rolling_mean sensor_measurement_2_rolling_std'
'sensor_measurement_2_rolling_std^2']

```



```

-----
-----
ValueError                                Traceback (most recent call
last)
Cell In[39], line 55

```

```

52 top_features_indices = [i for i in top_features_indices if i <
len(feature_names)]
54 # Plot Partial Dependence for the top valid features using
indices
--> 55 PartialDependenceDisplay.from_estimator(rf_poly_model,
X_train, top_features_indices, grid_resolution=50)
56 plt.show()

```

File ~\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\inspection_plot\partial_dependence.py:685, in PartialDependenceDisplay.from_estimator(cls, estimator, X, features, sample_weight, categorical_features, feature_names, target, response_method, n_cols, grid_resolution, percentiles, method, n_jobs, verbose, line_kw, ice_lines_kw, pd_line_kw, contour_kw, ax, kind, centered, subsample, random_state)

```

683 for i in chain.from_iterable(features):
684     if i >= len(feature_names):
--> 685         raise ValueError(
686             "All entries of features must be less than "
687             "len(feature_names) = {0}, got
{1}.".format(len(feature_names), i)
688         )
690 if isinstance(subsample, numbers.Integral):
691     if subsample <= 0:

```

ValueError: All entries of features must be less than len(feature_names) = 54, got 227.

```
import matplotlib.pyplot as plt
```

```
# Residuals
```

```
residuals = y_test - y_test_pred
```

```

plt.figure(figsize=(10, 6))
plt.scatter(y_test_pred, residuals, alpha=0.7)
plt.hlines(y=0, xmin=min(y_test_pred), xmax=max(y_test_pred),
color='red')
plt.xlabel('Predicted RUL')
plt.ylabel('Residuals')
plt.title('Residual Analysis')
plt.savefig('Residual Analysis.png')
plt.show()

```

